

1. Project Overview

You are tasked with building an agentic AI system in the form of a FastAPI microservice. This system will process free-form IT requests, determine which specialized agents are required, execute those agents to produce outputs, and return a comprehensive response that combines all results.

Core Requirements

Your system must:

1. Plan which specialist agents are needed based on the input request
2. Execute those LLM-powered agents to produce concrete outputs
3. Support approval workflows when requested by the caller
4. Return a structured JSON or Markdown response with combined results

No private knowledge base or database is required—your agents can rely on the capabilities of the LLM and/or call public data sources that you choose to integrate.

2. Technical Architecture

Your implementation must include the following components:

API Layer

FastAPI service with these endpoints:

- **POST /api/v1/execute** - Process new requests
- **POST /api/v1/plans/{id}/approve** - Approve a pending plan
- **POST /api/v1/plans/{id}/reject** - Reject a pending plan
- **GET /api/v1/tasks/{id}** - Get task status and results

Agent Layer

At minimum, you must implement these agents:

- **CoordinatorAgent**: Parses the request, creates execution plans, and merges outputs
- **DiagnosticAgent**: Performs root-cause analysis and provides ranked fixes
- **AutomationAgent**: Generates and syntax-checks PowerShell/Bash/Azure CLI scripts
- **WriterAgent**: Transforms results into structured content (email, SOP, summary)

You're welcome to add additional agents if you think they would enhance your solution.

Workflow Layer

You must use LangGraph for orchestrating the workflow:

- Implement a **CoordinatorGraph** for the overall execution flow
- Create at least one specialist graph with conditional edges (e.g., **DiagnosticGraph**)

LLM Integration

- Use OpenAI's or Any LLM

Advanced Techniques

Include at least one integration with DSPy, MCP such as:

- A DSPy router component
- An MCP context pruner
- Another appropriate integration of your choice

Testing

Implement the 5 specific pytest cases outlined in section 6

Documentation

- Complete README.md with setup and usage instructions
- Architecture diagram (using Mermaid, PNG,...)

3. Required Test Cases

You must implement these specific pytest cases:

Test Name	Purpose
test_happy_path	Verify that Example A returns "completed" status and includes non-empty diagnosis and script sections
test_approval_flow	Verify that Example B begins with "waiting_approval" and transitions to "completed" after approval

test_agent_retry	Simulate a failure in the AutomationAgent and verify that the Coordinator either retries or handles the failure gracefully
test_script_compiles	Verify that generated PowerShell/Bash scripts pass syntax checks (using <code>pwsh -Command {...}</code> or <code>bash -n</code>)

4. API Contract & Examples

4.1 POST /api/v1/execute

Request Fields:

Field	Type	Description
request	string	Natural-language description of what the assistant should do
require_approval	boolean	If true, pause after planning for explicit approval

Example A - Direct Execution (No Approval)

Request:

```

http
POST /api/v1/execute
Content-Type: application/json

{
  "request": "Diagnose why Windows Server 2019 VM cpu01 hits 95%+
CPU, generate a PowerShell script to collect perfmon logs, and draft
an email to management summarising findings.",
  "require_approval": false
}
```

Internal Flow:

Phase	Action
-------	--------

Plan CoordinatorAgent detects three sub-tasks (diagnostic, script, email)

Execute DiagnosticAgent → root cause & fixes
 AutomationAgent → PowerShell script
 + lint pass
 WriterAgent → management email

Merge Coordinator combines results into single response

Expected Response:

```
json
{
  "task_id": "123e4567-...",
  "status": "completed",
  "diagnosis": {
    "root_cause": "Wsappx runaway process",
    "evidence": ["perfmon shows high kernel time ..."],
    "solutions": [
      { "title": "Disable Superfetch", "confidence": "high" },
      { "title": "Install KB500XYZ", "confidence": "medium" }
    ]
  },
  "script": {
    "language": "powershell",
    "code": "New-Item C:\\logs -Force; logman start ...",
    "lint_passed": true
  },
  "email_draft": "Hello team, ...",
  "duration_seconds": 42
}
```

Example B - Approval Flow

Initial Request:

```
http
POST /api/v1/execute
Content-Type: application/json

{
  "request": "Create Azure CLI commands to lock RDP (3389) on my
three production VMs to 10.0.0.0/24 and pause for approval before
outputting the commands.",
  "require_approval": true
}
```

Initial Response (waiting for approval):

```
json
{
  "task_id": "plan-456",
  "status": "waiting_approval",
  "plan": {
    "steps": [
      "Generate NSG rules",
      "Generate rollback script"
    ],
    "summary": "Will restrict 3389 inbound to 10.0.0.0/24 on vm-a,
vm-b, vm-c"
  }
}
```

Approval Actions:

```
bash
# approve and resume
curl -X POST /api/v1/plans/plan-456/approve

# reject
curl -X POST /api/v1/plans/plan-456/reject
```

Response After Approval:

```
json
{
  "status": "completed",
  "commands": [
    "az network nsg rule create ...",
    "az network nsg rule delete ... # rollback"
  ]
}
```

4.2 GET /api/v1/tasks/{id}

Returns the current state of a task with one of these status values:

- active
- waiting_approval
- completed

- failed
- etc.

The response includes any partial results or error information.

5. Deliverables

Your submission must include:

1. Git repository (or ZIP file) containing all code, tests, README, and architecture diagram
2. A screen capture demonstrating Example A and Example B flows
3. Postman or curl collection for testing the API